



# SQL Indexes

Andreas Finger  
@mediafinger

# SQL

- Structured Query Language
- used by RDBMS (Relational DataBase Management Systems)  
e.g. MySQL, PostgreSQL, Oracle
- not used by noSQL databases ;-)  
e.g. Redis, CouchDB, MongoDB

# SQL

- In Rails we can formulate most queries with ActiveRecord
- but sometimes you will need SQL in Rails
- SQL can be much(!) faster than ActiveRecord
- and often you want to quickly get some data out of your database  
e.g. run some statistics, export some tables


# SQL

- Task.where(status: "todo", user\_id: 1)  
SELECT \* FROM tasks  
WHERE status = 'todo'  
AND user\_id = 1
- Task.where(status: ["todo", "doing"], user\_id: [1, 2]).  
.order(updated\_at: :asc)  
SELECT \* FROM tasks  
WHERE status IN ('todo', 'doing')  
AND user\_id IN (1, 2, 3, 4)  
ORDER BY updated\_at ASC


# SQL

- `Project.first.users`  
`SELECT users.* FROM users`  
`INNER JOIN projects_users ON users.id =`  
`projects_users.user_id`  
`WHERE projects_users.project_id = 1`
- `User.where(confirmed: true).`  
`joins(:tasks).where(:tasks, status: ["todo", "doing"])`  
`SELECT users.* FROM users`  
`INNER JOIN tasks ON tasks.user_id = users.id`  
`WHERE users.confirmed = 1`  
`AND tasks.status IN ('todo', 'doing')`


# SQL JOINS




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
INNER JOIN TableB B  
ON A.Key = B.Key
```




```
SELECT <select_list>  
FROM TableA A  
LEFT JOIN TableB B  
ON A.Key = B.Key  
WHERE B.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
RIGHT JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key
```



```
SELECT <select_list>  
FROM TableA A  
FULL OUTER JOIN TableB B  
ON A.Key = B.Key  
WHERE A.Key IS NULL
```

# SQL Indexes

- You should create an Index, when you
  - access elements on a table often with foreign keys  
e.g. `Tasks.where(user_id: 1)`
  - often use queries that join tables
  - sort by a large dataset by a column often

# SQL Indexes

- When you create an Index, SQL needs some time to go through the data, sort it and write the Index to disk
- That means every Index costs disk space
- Indexes that cover several fields, are harder to create the right way, as the order of the fields matter
- You only need an Index if you have A LOT of data as RDBMS are made to handle many millions of entries



# SQL Indexes

- Indexes can speed up queries a lot
- For large datasets and complex operations it could look like this:

30 seconds	(no index)
4 seconds	(a not perfect index)
0.08 seconds	(the right index)
- any normal query should just take a few milliseconds



- Databases are optimized to search, join, aggregate, group and order data
- They can also do calculations much faster than it could be done on the server side
- Whenever possible, let your database do the heavy lifting and just work with the results

# noSQL

- Key-Value Stores are a type of noSQL databases
- Key-Value-Stores are often used for Caching, or to save the messages of a Queue
- Key-Value Stores are very fast, because they are very simple
- A popular Key-Value Store is Redis
- Postgres has a Key-Value Store build in: HSTORE

# noSQL

- Document Stores are a type of noSQL databases
- DS excel at saving complete datasets without fixed form or size and allowing to access them
- DS usually have good full text search capabilities
- DS have are not great when running reporting over the whole database
- the flexibility of DS looks great at the beginning, but can become technical debt later